

pif_eth_100 — User's Guide: run the 100 Mbaud TX and RX in the browser UI

PRU0 sends 200-byte 8b/10b BERT frames at 100 Mbaud (300 MHz core, TX divider 3); PRU1 receives them through a loopback wire at 2x oversampling (divider 1.5) and checks CRC-32 and bit errors against the same PRNG. When a frame is done DRAM1 reads `crc_ok 1`, `symbol_errors 0`, `bit_err 0`, `tot_bits 1600`. Simulator only, never run on silicon.




Figure 1. What the two cores do: PRU0 (steps 1 to 5) sends, PRU1 (steps 6 to 10) receives and writes the DRAM1 stats you read in the steps below. Simulator model, not silicon.

How this guide was checked. UI wording in this guide was derived from the UI source and validated by scripted walkthrough over the same HTTP/WebSocket endpoints; a real-browser click-through was not performed in the authoring session.

What goodput to expect. Inside a burst the wire runs at 100 Mbaud = 80 Mbit/s, but a frame is only sent after PRU1 has finished decoding the previous one, so the end-to-end goodput is the payload divided by the whole frame period.




Figure 2. One frame period for the baseline RX, the fast RX and TX alone (simulator run via `run_100`, not silicon). Regenerate the figures with `python3 source/pif_eth_100/make_figures.py`.

Prerequisites

- `pip install -r requirements.txt` (the seed/status client needs `websockets`).
- Run every command from the repo root.

Steps

1. **Config: 300 MHz on both cores.** The server always loads the repo-root `memory.cfg`:

```
cp memory.cfg memory.cfg.bak && cp config/memory_pif_eth_100.cfg memory.cfg
python3 ui/server.py
```

8080 is the server's default port, and the rest of this guide uses it. (The scripted walkthrough that validated the guide used a private port, 8091, via `python3 ui/server.py --port 8091`.) `memory.cfg` is a tracked file; step 9 restores it.

Alternative for a server that is already running: pick **300 MHz** in the PRU speed dropdown in the controls bar **before doing anything else**. The dropdown rewrites `memory.cfg` and rebuilds the simulator, so it drops loaded programs, the loopback and all DRAM contents. You will need to restore `memory.cfg` yourself afterwards (`git checkout -- memory.cfg`).

2. **Open the UI** at `http://localhost:8080` and confirm the speed dropdown in the controls bar reads **300 MHz**.
3. **Multi-core on, partner PRU1.** Click **Multi-core** in the controls bar, then choose **PRU1** in the partner dropdown that appears next to it (its default is RTU0). Multi-

core Run steps PRU0 as the lead and paces PRU1.

4. **Load both programs from the in-UI source browser.** Click **Project** (in the Assembly Editor panel title bar) and choose `pif_eth_100`. The folder's files open as tabs: `pif_eth_100_rx.asm`, `pif_eth_100_rx_fast.asm` (optional, step 4c), `pif_eth_100_tx.asm`, `pif_eth_crc32_hw.inc`.

- Select the `pif_eth_100_tx.asm` tab, set the load-target dropdown in the editor title bar (it appears once Multi-core is on) to → **PRU0**, click **Load & Assemble**.
- Select the `pif_eth_100_rx.asm` tab, set the load target to → **PRU1**, click **Load & Assemble**.

Warnings:

- Do not use **Open** → "Browse file system..." (the OS file picker). It drops the folder from the file path, and `.include "pif_eth_crc32_hw.inc"` then cannot be found. (The **Open** list itself shows only top-level `source/` files, not the `pif_eth_100` folder; use **Project**.)
- Never load the same file into both cores.

4b. **Click Reset.** In multi-core mode it resets PRU0 and the partner core (PC and registers to 0). **Load does not reset a core's PC or registers.** Loading `pif_eth_100_rx_fast.asm` over a previously run `pif_eth_100_rx.asm` without Reset resumed at the old PC and wrote LUT bytes over the stats block (reproduced by the planner; not re-run for this guide). Always Reset after loading, then seed.

4c. **Optional: the optimised receiver.** To run the optimised receiver, load `pif_eth_100_rx_fast.asm` into PRU1 instead of `pif_eth_100_rx.asm` (select its tab in the **Project** list, load target → **PRU1**, **Load & Assemble**), then **Reset** as in step 4b. Seeding and status are identical (steps 5 to 8). To switch between the two receivers, load the other file and Reset again. The optimised receiver finishes decoding sooner, so a frame is done after fewer lead instructions than with the baseline (see the README for figures).

5. **Seed.**

```
python3 source/pif_eth_100/seed_ui_100.py
```

Expected output:

```
seeded DRAM0 (TX) + DRAM1 (RX, RXCFG 0x0000801F), loopback ch0 latency 0.8333 ns, frame 1 armed -- click Run
```

This writes PRU0's TX control block and 8b/10b LUT (DRAM0), PRU1's decode LUT and control block (DRAM1), zeroes the stats block, enables **loopback channel 0 with latency 5/6 ns = 0.8333 ns** (`run_100.LOOPBACK_LATENCY_NS`), and arms frame 1 (TX go flag, then RX go flag). The client always talks to PRU0's loopback, which is where the UI's loopback lives.

Where to look: the Loopback card is in the I/O panel's Peripheral Interface section, titled "Loopback → PRU1 RX (jitter / latency / drift)", with one row per channel (`ch0` checkbox, `lat ns`, `jit ns`, `drift ppm`, **Apply**). That section is shown once the firmware has switched the pin mux to the Peripheral Interface, so it may only appear after the first Run.

- Keep the current core PRU0 if you press **Apply**; it applies to the current core.
- The `lat ns` box has a spinner step of 0.5. The server reports `0.8333333333333334`; leave the box untouched and Apply keeps that value. If you type it, `0.8333` is fine (see "Loopback latency" below).

(The two points above are from the UI source and the server's state message (`io.loopback` was read back as `ch0` enabled, latency `0.8333333333333334 ns`, jitter 0, drift 0). They were not observed in a browser; see Validated.)

6. **Click Run** (PRU0 as the current core). Multi-core Run steps 1000 PRU0 instructions every 10 ms and paces PRU1; it does not stop by itself. One frame needs at most 35 000 PRU0 instructions (the scripted walkthrough checks after every 5000 lead instructions and saw `frames` become 1 by the 35 000-instruction check), so click **Stop** after a few seconds, once `status` (step 7) shows `frames 1`. Optionally, in the Peripheral panel or a Memory read, check the perif config registers: PRU0 TXCFG (`0x260E4`) = `0x00020010`, PRU1 RXCFG (`0x26100`) = `0x0000801F`. Both read back exactly these values here.

7. Check the result.

```
python3 source/pif_eth_100/seed_ui_100.py status
```

Observed (frame 1 complete), exit code 0:

```
frames          1
cap_bytes       526
rx_ovf          0
symbol_errors   0
crc_ok          1
bit_err         0
tot_bits        1600
eof_status      1
PASS
```

Or use a Memory panel at `0x2F00` length 64 for the stats and `0x2E00` length 204 for the reconstructed frame (200 B payload + 4 B FCS). The payload is the xorshift32 stream of the seed, continuing across frames: frame 1 starts `a521c86ebe053c0c` (golden model `prng_bytes(8, 0x1BADC0DE)`), frame 2 starts `b0a1abf161f49863`. The second was read back with `read_memory` at `0x2E00` after the 2-frame scripted walkthrough on 2026-09-30 (FCS bytes `ed97f06d`); the first is the model's value, not read from the UI.

`status` exits 1 and prints `FAIL (...)` until `frames` reaches 1. While RX is still decoding (`post_frame`), an intermediate state (for example `frames 0` and `crc_ok 0`

with a non-zero `cap_bytes`) is reported as FAIL. That is not an error; Run a little longer and query again.

8. **Next frame.** `python3 source/pif_eth_100/seed_ui_100.py arm`, then click **Run** again (about the same instruction count). `status` then shows `frames 2`, `crc_ok 1`. Each frame needs its own `arm`.

9. **Restore.** Stop the server, then `mv memory.cfg.bak memory.cfg` (or `git checkout -- memory.cfg`; then `git status` must not list it).

DRAM1 (PRU1 RX) addresses

Host (global) addresses; the firmware sees them minus `0x2000`.

Address	Contents
0x2000	8b/10b decode LUT, 1024 x u16 (2048 B)
0x2800	raw 2x-oversampled capture, 1024 B (about 528 B used; <code>cap_bytes</code> 526)
0x2C00	unused by the baseline RX (the fast RX builds its 256 B decimation LUT here)
0x2E00	reconstructed frame, 256 B (200 B payload + 4 B FCS)
0x2F00	stats <code>frames</code> (u32)
0x2F04	stats <code>cap_bytes</code>
0x2F08	stats <code>rx_ovf</code> — nonzero = RX FIFO overflowed
0x2F0C	stats <code>symbol_errors</code> — 8b/10b decode failures
0x2F10	stats <code>crc_ok</code> — 1 = FCS matched
0x2F14	stats <code>prng_bit_errors</code> (<code>bit_err</code>)
0x2F18	stats <code>tot_bits</code> (= 1600 for a 200 B frame)
0x2F1C	stats <code>eof_status</code> — 1 = clean EOF
0x2F40	control <code>mode</code> (0 = BERT)
0x2F44	control <code>seed</code> (must match PRU0's)
0x2F48	control <code>payload_len</code> (200)
0x2F4C	control <code>go</code> (host arms per frame)
0x2F50	control <code>rxcfg</code> (0x0000801F)

The TX control block is in DRAM0: `0x0400` `num_frames`, `0x0404` `mode`, `0x0408` `seed`, `0x040C` `payload_len`, `0x0410` `frame count`, `0x0414` `burst bytes pushed`, `0x0418` `go flag`.

Troubleshooting

Symptom	Cause	Fix
empty capture, <code>frames</code> stays 0	loopback off, or TX loaded into both cores	re-run <code>seed</code> , check which file each core has
<code>crc_ok 0</code> / <code>eof_status 0</code> / garbage	seeded after PRU1 had already run (control block read at boot)	reload PRU1, re-seed, then Run

Symptom	Cause	Fix
programs vanished	speed dropdown used after loading	set 300 MHz first, then reload
stats full of garbage right after Run (e.g. frames 185272842)	a core was re-loaded without Reset and resumed at its old PC	Reset both cores, re-seed, Run
.include not found on Load	file opened with the OS file picker	open it from the in-UI Project browser
seed_ui_100.py says "needs 300 MHz"	server is on another clock	step 1

The rows above come from the plan and the seed/status client's messages. The "garbage after Run without Reset" row was reproduced by the planner, not re-run here. The frames 0 row (no seed yet) was observed: status before Run, right after seeding, printed all zeros and FAIL (frames, crc_ok, tot_bits, eof_status), exit 1.

Notes:

- The Signal Graph only traces the lead core (PRU0) during a multi-core Run. To watch RX state live, switch the current-core selector to PRU1 (or leave multi-core) and read the PRU1 Peripheral Interface channel card (from the UI source, not clicked; see the notice at the top).
- **Loopback latency and the known RX limit.** RX decode is not robust when a sample lands exactly on a TX edge (0 ns wire delay at one of the three sample-to-edge phases): test_all_edge_phases_clean at the 10/3 ns shift failed with 71 symbol errors at 0 ns wire delay (reproduction command in the README "Latency note"). This project's loopback therefore uses 5/6 ns = 0.8333 ns (run_100.LOOPBACK_LATENCY_NS), which keeps every sample at least 0.83 ns from a TX edge. seed_ui_100.py sets that value for you. In the UI walkthrough the arm phase is fixed by the host's paced stepping, so a single-phase UI run does not exercise every phase: with latency set to 0.0, 0.5, 0.8333, 5/6 and 1.0 ns, one frame each passed clean in this UI flow (see Validated). One frame passing in the UI arm phase is not evidence that 0 ns is robust; keep 0.8333 ns. The lat ns field has a 0.5 ns spinner step, so the value is not reachable by clicking the arrows; leave what seed set, or type 0.8333.

Validated

Date: 2026-09-30. Simulator only (not silicon). Server on a private port (8091) with config/memory_pif_eth_100.cfg as memory.cfg. All four scripted runs below (base, fast, base, fast) were repeated at 14:16 UTC at commit 6f4b458 with identical output, and memory.cfg was restored byte-identical afterwards.

- **Browser walkthrough not possible in this session** (the Chrome extension was not connected: tabs_context_mcp reported "Browser extension is not connected" on two attempts). The scripted walkthrough over the same endpoints is the validation of record. The browser-only wording in this guide (button labels, dropdown names, where the Loopback card is, the **Project** flow) comes from

reading `ui/static/index.html`, `app.js` and `server.py`, and the `/source` listing (which returned `pif_eth_100` with its four files), not from clicking in a browser.

- **Scripted walkthrough** (`python3 source/pif_eth_100/ui_walkthrough_100.py --frames 2 --ws ws://127.0.0.1:8091/ws --http http://127.0.0.1:8091`), exit 0:

```
[step 1-2] UI clock 300 MHz
[step 4] loaded pif_eth_100_tx.asm -> pru0
[step 4] loaded pif_eth_100_rx.asm -> pru1
[step 4b] Reset PRU0 + PRU1 (PC/registers to 0)
[step 5] seeded DRAM0/DRAM1, loopback ch0 on (latency 0.8333 ns), frame 1 armed
[step 6-7] frame 1 after 35000 lead instr: {'frames': 1, 'cap_bytes': 526, 'rx_ovf': 0, 'symbol_errors': 0, 'crc_ok': 1, 'bit_err': 0, 'tot_bits': 1600, 'eof_status': 1} -> PASS
[step 8] armed next frame
[step 6-7] frame 2 after 35000 lead instr: {'frames': 2, 'cap_bytes': 526, 'rx_ovf': 0, 'symbol_errors': 0, 'crc_ok': 1, 'bit_err': 0, 'tot_bits': 1600, 'eof_status': 1} -> PASS
UI WALKTHROUGH PASS (2 frames, base RX)
```

- **Optimised receiver (step 4c), scripted only.** Same server setup (private port 8091, 300 MHz). `ui_walkthrough_100.py --frames 2 --rx fast`, then `--rx base`, then `--rx fast` again (each run loads both cores and Resets both, so this covers the fast -> base -> fast re-load with Reset), all exit 0. Each fast run: frame 1 frames 1, cap_bytes 526, rx_ovf 0, symbol_errors 0, crc_ok 1, bit_err 0, tot_bits 1600, eof_status 1 PASS after 20 000 lead instructions, frame 2 the same fields with frames 2 PASS after a further 15 000 (the base runs needed 35 000 each). This was validated by script and by the `_ui_flow_one_frame("fast")` unit test only; the browser flow for the optimised receiver (choosing the fast tab in **Project**, load target, **Reset**) was not clicked through in a browser. The four-tab **Project** listing is inferred from the server's directory listing, not observed in a browser.
- **Seed/status CLI flow (steps 5, 7, 8).** The scripted walkthrough calls the same functions as `seed_ui_100.py` (seed, arm, status). A hand-typed CLI session was also tried during authoring, but its instruction counts were not saved and are not quoted here.
- **Loopback latency probe** (fresh load, Reset and seed each time, latency then overridden, 60 000 lead instructions, one frame each): 0.0, 0.5, 0.8333, 0.8333333333333334 and 1.0 ns all gave frames 1, crc_ok 1, symbol_errors 0, bit_err 0, tot_bits 1600, eof_status 1, rx_ovf 0.
- Read back from the server after a run: PRU0 TXCFG 0x00020010, PRU1 RXCFG 0x0000801F, loopback ch0 enabled, latency 0.8333333333333334 ns.